



BITTIGER

DS 501 Data scientist express bootcamp

Week 1 Knowledge

Copyright Policy

All content included on the Site or third-party platforms as part of the class, such as text, graphics, logos, button icons, images, audio clips, video clips, live streams, digital downloads, data compilations, and software, is the property of BitTiger or its content suppliers and protected by copyright laws.

Any attempt to redistribute or resell BitTiger content will result in the appropriate legal action being taken.



We thank you in advance for respecting our copyrighted content.

For more info:

see <https://www.bittiger.io/termsfuse>

and <https://www.bittiger.io/termservice>

版权声明

所有太阁官方网站以及在第三方平台课程中所产生的课程内容，如文本，图形，徽标，按钮图标，图像，音频剪辑，视频剪辑，直播流，数字下载，数据编辑和软件均属于太阁所有并受版权法保护。

对于任何尝试散播或转售BitTiger的所属资料的行为，太阁将采取适当的法律行动。



有关详情，请参阅

<https://www.bittiger.io/termsfuse> <https://www.bittiger.io/termservice>



What is statistics? Why is it important?

BIT TIGER



Statistics

- All about quantifying uncertainty (which subject is studying certainty?)
- Why uncertainty? **imperfect** data
 - Psychology
 - Geology
 - Economics
- Why imperfect data?
 - Only observe a small fraction, e.g., candidate poll
 - Only observe indirect signs, e.g., neuro signal
 - Data always contain noise, e.g, measurement error

<http://www.stat.cmu.edu/~cshalizi/uADA/16/lectures/01.pdf>



Probability

- Probability is central to statistics
 - Probability = Statistics?
- Basic concepts
 - **Experiment**
 - **Sample space:** all possible outcomes of an **experiment**
 - **Event:** subset of sample space

<https://www.youtube.com/watch?v=KbB0FjPg0mw&list=PL2SOU6wwxB0uwwH80KTQ6ht66KWxbzTlo>
from 15min

Statistics uses probability to model inference from data. We try to mathematically understand the properties of different procedures for drawing inferences: Under what conditions are they reliable? What sorts of errors do they make, and how often? What can they tell us when they work? What are signs that something has gone wrong? L



Random variable and types

- Random variable (X)
 - Enumerate sample space to real value
- Categories
 - Discrete
 - Continuous
- Examples
 - Flip a coin; Daily traffic; Number of people click on an ad; BMI
- R
 - bar chart, pie chart
 - Probability density



How to describe probability of random variable

- Probability density function (PDF)



- PDF: $x \sim P(X = x)$
- For X , frequency of possible outcome x 's occurrences of the total possible occurrences

■ Example: rolling dice, $P(X = 1) = \frac{1}{6}$

- Discrete (pmf) v.s continuous (pdf)
- Non-negative, sum up to 1

- Rules

- $0 \leq P(X = x) \leq 1$, $P(X = \Omega) = 1$, $P(X = \varnothing) = 0$, $P(X = x) + P(X \neq x) = 1$,
 $P(X = x \text{ or } Y = y) = P(X = x) + P(Y = y) - P(X = x \text{ and } Y = y)$





Example interview questions

- 2 dices, probability of at least one 4 shows up
- 2 dices, probability of the sum is 6
- ...

BITTIGER






How to describe probability of random variable(cont'd)

- Cumulative density function (CDF)


- CDF: $x \sim P(X \leq x)$
- Discrete v.s continuous
- Non-decreasing, up to 1

- Examples

- Roll dice
- Uniform distribution



Relations Between Two Different Typical Representations of a Population






Conditional probability

- How to update uncertainty after getting new information?
 - $P(\text{dice} = 1) = \frac{1}{6}$, what if we know it's an odd number?
- Definition, Probability of A given B

$$P(A | B) = \frac{P(A \cap B)}{P(B)}$$

- Intuition
 - Only consider B
 - Renormalize
- Example question





Independence

- Def: $P(A \cap B) = P(A) \times P(B)$
- Considering conditional probability $P(A | B) = \frac{P(A \cap B)}{P(B)}$

$$P(A | B) = \frac{P(A)P(B)}{P(B)} = P(A)$$

- A and B are conditionally independent given C

- $P(AB | C) = P(A|C) P(B|C)$

conditional independence is the same as normal independence, but restricted to the case where you know that a certain condition is or isn't fulfilled.

Say you roll a blue die and a red die. the results of the dice rolls are conditionally independent with respect to the event "the blue result is not 6 and the red result is not 1", but they're not conditionally independent with respect to the event "the sum of the results is even".



A and B are independent \leftrightarrow
A and B are independent given C ?

Left to right. Independence does not imply conditional independence: for instance, independent random variables are rarely independent conditionally on their sum or on their maximum.

Right to left: $A = x + C$, $B = y + C$

$$P(AB) = P(A)P(B)$$

$$C = A - B, \text{ then } P(AB|C) \neq P(A|C) P(B|C)$$

$$A = C + D, B = C + E, \text{ then } P(AB|C) = P(A|C) P(B|C) \text{ but } P(AB) \neq P(A) P(B)$$



Bayes' theorem

- Reverse the conditional probability

$$P(B | A) = \frac{P(A | B)P(B)}{P(A | B)P(B) + P(A | B^c)P(B^c)}$$

- Prove it

○ Hint:

$$P(A | B) = \frac{P(A \cap B)}{P(B)}$$

- Fair coin and unfair coin problem

$P(\text{fair coin} | 8 \text{ heads out } 10) = P(8 \text{ heads out } 10 | \text{fair coin}) * P(\text{fair coin}) / [P(8 \text{ heads out } 10 | \text{fair coin}) * P(\text{fair coin}) + P(8 \text{ heads out } 10 | \text{not fair coin}) * P(\text{not fair coin})]$

B: Gene mutation A: medical test

http://web.ipac.caltech.edu/staff/fmasci/home/astro_refs/Science-2013-Efron.pdf



Expectation

- Measure what's the center of a random variable
 - Discrete $E[X] = \sum_x xp(x)$
 - Continuous $E(X) = \int_x xP(X=x)$
 - Roll dice, flip coin, ... distribution
- Population mean v.s. sample mean
 - Sample mean **itself** is random variable, estimator of population mean
 - Unbiased: $E(\text{sample mean}) = \text{population mean}$
- Why always larger sample is better?
 - Biased when fewer data?
 - Proof. Hint: $E(X + Y) = E(X) + E(Y)$

一个人去赌场，花5刀玩一个游戏，扔两次骰子，如果和为6，他赢21刀，其他就什么也没有。问，这个游戏是偏向casino还是player的。（就是算个期望，然后期望是负的所以是偏向casino）

follow up: 如果现在有一个策略，这个人一直玩，直到第一次赢，然后走，问，他赢的概率是多少。

<https://stats.stackexchange.com/questions/30365/why-is-expectation-the-same-as-the-arithmetic-mean>



Variance

- Variance (σ^2) measure how random variable spreads

$$Var(X) = E[(X - \mu)^2] = E[X^2] - E[X]^2$$

- Standard deviation (σ)
- Sample variance (s^2) and standard deviation (s)

$$s^2 = \frac{\sum_{i=1}^n (X_i - \bar{X})^2}{n - 1}$$

- Why (n-1): Unbiased estimator of σ^2
- Proof (HW), Hint $Var(X+Y) = Var(X) + Var(Y)$ if X, Y independent

$$E(X+Y) = E(X) + E(Y); E(k*X) = k*E(X)$$

$$E(S^2) = Var(X)$$



Important discrete distribution

- Bernoulli

$$P(X = x) = p^x (1 - p)^{1-x}$$

- Binomial

- Sum of N Bernoulli variables

-

$$P(X = x) \equiv \binom{N}{x} p^x (1 - p)^{N-x}$$

$$E(X) = p, \text{Var}(X) = p(1-p)$$



What's the mean and variance?

$$E(X) = N \cdot p, \text{Var}(X) = N \cdot p(1-p)$$

Poisson: a given number of events occurring in a fixed interval of time



More details about other distributions

- Geometric distribution

- Prob of seeing first success with k independent trials, each with success probability p .

$$\Pr(X = k) = (1 - p)^{k-1} p$$

- Negative binomial distribution

- Prob of seeing r failures with k success trials, each with success probability p .

$$\Pr(X = k) = \binom{k+r-1}{k} p^k (1-p)^r$$



Important distribution

- Poisson, count data per time unit

$$P(X=k) = \frac{\lambda^k}{k!} e^{-\lambda}$$

○ examples

- Exponential distribution



○ In a poisson distribution, inter-arrival time follows exponential distribution

○ PDF: $\lambda e^{-\lambda x}$ CDF: $1 - e^{-\lambda x}$

○ Expectation $E[X] = \frac{1}{\lambda}$

○ Variance $Var[X] = \frac{1}{\lambda^2}$

○ "Memoryless"
 $P(X > s+t \mid X > s) = P(X > t)$



Proof of memoryless: <http://stats.stackexchange.com/questions/2092/relationship-between-poisson-and-exponential-distribution>

Really good read of binomial and poisson connection:

<https://probabilityandstats.wordpress.com/2011/08/18/poisson-as-a-limiting-case-of-binomial-distribution/>

lambda is the mean arrival rate per unit time

$1 - e^{-(\lambda * x)}$



Important continuous distribution

- Normal distribution

$$P(X = x) = (2\pi\sigma^2)^{-1/2} e^{-(x-\mu)^2/2\sigma^2}$$

$$E[X] = \mu \text{ and } \text{Var}(X) = \sigma^2$$


- Standard normal distribution (Z)

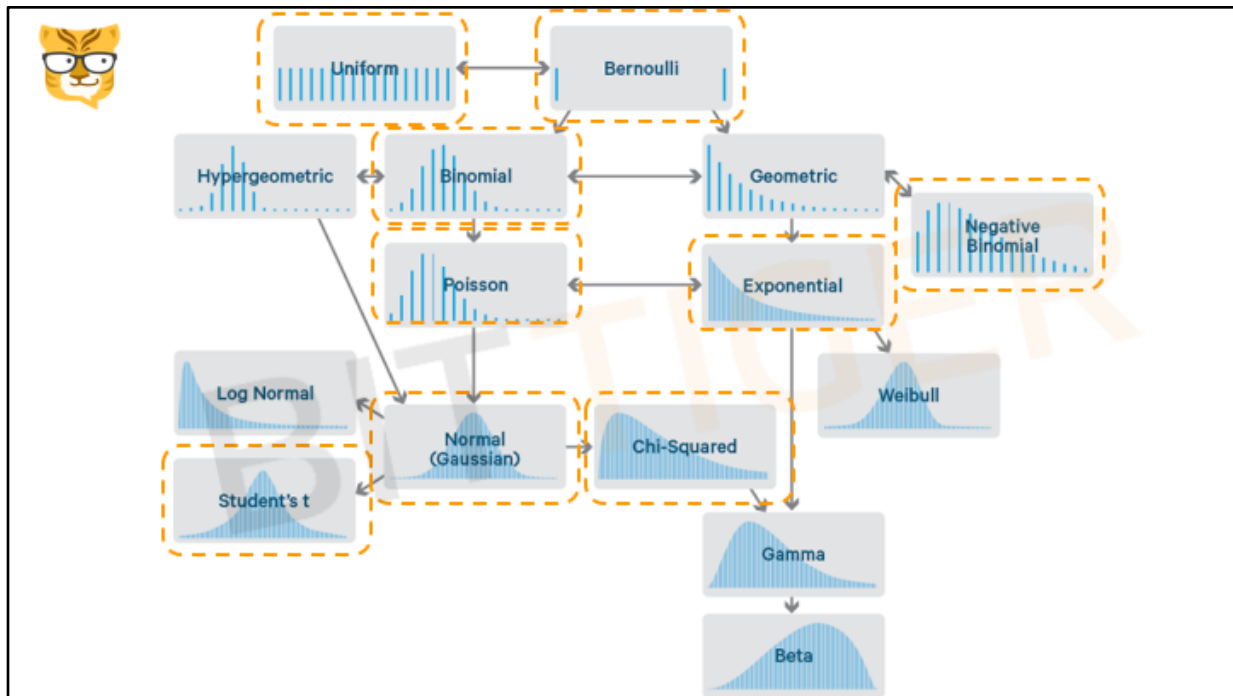
$$\mu = 0 \text{ and } \sigma = 1$$

$$X = \mu + \sigma Z \sim N(\mu, \sigma^2)$$

- beauty of normal curve (6σ)

- $[\mu - 3\sigma, \mu + 3\sigma]$ covers 99.7%
- $[\mu - 2\sigma, \mu + 2\sigma]$ covers 95%
- $[\mu - \sigma, \mu + \sigma]$ covers 68%





<http://blog.cloudera.com/blog/2015/12/common-probability-distributions-the-data-scientists-crib-sheet/>
<http://individual.utoronto.ca/zheli/poisson.pdf>



Example interview questions

- In a country in which people only want boys every family continues to have children until they have a boy. If they have a girl, they have another child. If they have a boy, they stop. What is the proportion of boys to girls in the country?
 - What distribution is this?
 - What's expected number of kids (X) each family has?
 - What the ratio of boys to girls?

<http://mathoverflow.net/questions/17960/google-question-in-a-country-in-which-people-only-want-boys>

(Geometric) $P(X = k) = p(1-p)^{k-1} = (1/2)^k$, where k is the number of children a family has until they have first boy

$$E(X) = \text{Sum}(k \cdot P(X=k)) = 1 \cdot \frac{1}{2} + 2 \cdot \left(\frac{1}{2}\right)^2 + \dots$$

$$E(X) \cdot \frac{1}{2} = 1 \cdot \left(\frac{1}{2}\right)^2 + 2 \cdot \left(\frac{1}{2}\right)^3 + \dots$$

$$\text{So } E(X) = 2$$

Then ratio is 50%



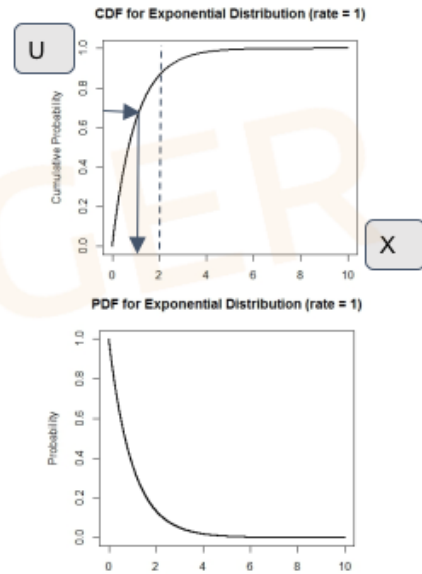
Another type of distribution Interview question

- How can we generate certain distribution?
 - Q1: Generate exponential distribution from uniform distribution
 - Q2: Generate normal distribution from uniform distribution



Inverse transform sampling

- To generate X with PDF: $P(X)$
 - Calculate CDF of X : $F(X)$ and inverse $F^{-1}(\cdot)$
 - Generate U with uniform distribution, say u
 - Calculate $F^{-1}(u)$
 - Repeat n times
- Intuition
- Let's solve Q1
- Can we solve Q2 as well?



Need to prove generated $X = F^{-1}(u)$ follows required PDF: P

Meaning we need to prove X follows required CDF: F ,
Let's see

$$\text{CDF}(X = x) = \text{Prob}(X \leq x) = \text{Prob}(F^{-1}(u) \leq x) = \text{Prob}(u \leq F(x)) = F(x)$$


https://en.wikipedia.org/wiki/Inverse_transform_sampling

<http://web.ics.purdue.edu/~hwan/IE680/Lectures/Chap08Slides.pdf>



Acceptance rejection sampling

- Simple version: if good, accept; else, reject.
 - Ex1: Generate Y with uniform $[\frac{1}{4}, 1]$ given $X \sim [0, 1]$
 - Ex2: Generate random integer Y in $[1, 7]$ given random integer X in $[1, 5]$ $5*(X-1) + X$
 - Ex3: Generate normal distribution given uniform distribution
 - Sample a point, x, calculate $P(x)$
 - Generate random number u from uniform distribution
 - If $u < P(x)$, accept x, else reject, back to step 1
- Check out this [video](#)



https://en.wikipedia.org/wiki/Inverse_transform_sampling

$(\text{rand5}() - 1) * 5 + \text{rand5}()$ gives us $[1, 25]$, only accept $[1, 21]$ then mod 7.

<http://www.columbia.edu/~ks20/4703-Sigman/4703-07-Notes-ARM.pdf>

http://www.tc.umn.edu/~elock/Rejection_Sampling.pdf

<https://www.youtube.com/watch?v=tiWQwllSOuo>

Picture from: http://people.duke.edu/~ccc14/sta-663-2016/15A_RandomNumbers.html



Percentile

- Definition (c%, N)



- Special percentiles
 - Median
 - 1st quantile
 - 3rd quantile
- Median, quantiles in normal distribution?
 - Look up **z score**, alpha: area outside $[-z, z]$



Why normal distribution is so important?



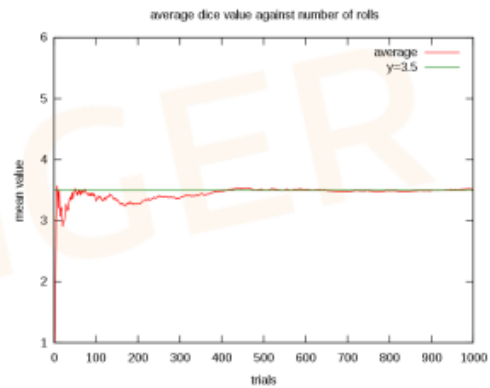
Law of large numbers

$X_1, X_2, \dots, X_n, \dots$ are independent, identically distributed (IID) random variables, X has finite mean μ .

- Sample mean converge to $\bar{X}_n \equiv \frac{1}{n} \sum_{i=1}^n X_i$

as n increases:

$$\bar{X}_n \xrightarrow[n \rightarrow \infty]{} \mu$$



Extended material: <http://www.math.uah.edu/stat/sample/LLN.html>



Central limit theorem

- $X_1, X_2, \dots, X_n, \dots$ are independent, identically - distributed (IID) random variables, X has finite mean μ and variance σ^2


(re $\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i \sim N(\mu, \sigma^2/n)$ standard error, CLT still holds)

- Application

- Binomial distribution

- Flip coin example:

$$\bar{p} \sim N\left(p, \frac{p \cdot (1-p)}{n}\right)$$



The central limit theorem is a very important tool for thinking about sampling distributions – it tells us the shape (normal) of the sampling distribution, along with its center (mean) and spread (standard error)



Confidence intervals

- Recall normal curve
 - What range covers 95% of possibility? 99%?
- Confidence interval (CI)
 - Derive 95% CI of observed \bar{p} in Binomial distribution
$$p \pm z_{1-\alpha/2} \sqrt{\frac{p \cdot (1-p)}{n}}$$
$$\hat{p} \pm z_{1-\alpha/2} \sqrt{\frac{\hat{p} \cdot (1-\hat{p})}{n}} \quad \hat{p} \pm \frac{1}{\sqrt{n}} \text{ approx.}$$
 - How many voters in the candidate poll?
 - Can Trump relax? How many are enough?



Fair coin $p = \frac{1}{2}$, what's 95% CI of **observed** p ?

What if p is not known, use estimator to replace



When CLT doesn't hold?

- Normal distribution -> T distribution when $N < 30$


- T distribution has only one parameter: degree of freedom ($df = N-1$)
- Approximate normal as df increases
- CI under normal distribution

$$Mean_{estimate} \pm z_{1-\alpha/2} * StdErr_{estimate}$$

- CI under t distribution

$$Mean_{estimate} \pm t_{n-1} * StdErr_{estimate}$$

- z or t?




t distribution was specially designed to provide more conservative test results when analyzing small samples (such as in the brewing industry).



What about CI of not normal distribution?



Skewed distribution



- What caused skewness
 - Bounded: lower bound or upper bound
- Can we still use normal distribution to estimate mean of skewed variable?



How to treat skewed variable

- Log transformation
- Use median, if so, what's CI of median?
 - Non-parametric method (v.s. parametric method)
 - **Bootstrap**
 - Jackknife



Covariance and correlation

- Definition

- $cov(X, Y) = E[(X - E[X])(Y - E[Y])]$

$$\rho_{X,Y} = corr(X, Y) = \frac{cov(X, Y)}{\sigma_X \sigma_Y}$$

- What is sample covariance and correlation?

- $cov(X, Y) = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) = \frac{1}{n-1} (\sum x_i y_i - n \bar{x} \bar{y})$

$$cor(X, Y) = \frac{cov(X, Y)}{s_X s_Y}$$

If $X=Y$, $E[(X - E(X))^2]$



Properties of correlation


- Properties

- $\text{cor}(X, Y) = \text{cor}(Y, X)$
- $-1 \leq \text{cor}(X, Y) \leq 1$
- $\text{cor}(X, Y) = 1$ or -1 : X, Y aligned on a line perfectly
- $\text{cor}(X, Y) = 0$: X, Y are independent

- Geometric interpretation

$$\rho_{X,Y} = \frac{\text{cov}(X,Y)}{\sigma_X \sigma_Y} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}} = \frac{\sum_{i=1}^n x'_i y'_i}{\|x'\| \|y'\|} = \cos(\theta)$$

- $\text{cor}(X, Y) = 0, \text{cor}(Y, Z) = 0 \Rightarrow \text{cor}(X, Z) = 0?$



<http://stats.stackexchange.com/questions/97051/building-the-connection-between-cosine-similarity-and-correlation-in-r>